

Nanostructured Materials

Machine Learning Basics

Attila Cangi

Center for Advanced Systems Understanding,
Helmholtz-Zentrum Dresden-Rossendorf

a.cangi@hzdr.de

June 17, 2025

Lecture Schedule

Teaching team (Shirong/Ebru/Chenchen/Attila)				
Lecture No.	Module	Lecture Topic	Date (Tue.)	Lecturer
1	M1: Introduction	General Introduction	08.04	Dr. Shirong Huang/Dr. Ebru Cihan
2	M2: Fabrication and characterization	Nanofabrication: Top-down/Bottom-up approaches	15.04	Dr. Chenchen Wang
3	M2: Fabrication and characterization	Imaging/characterization techniques: Electron microscopy based techniques	22.04	Dr. Ebru Cihan
4	M2: Fabrication and characterization	Scanning probe microscopy (SPM) 1: Scanning tunneling microscopy (STM)	29.04	Dr. Ebru Cihan
5	M2: Fabrication and characterization	Scanning probe microscopy (SPM) 2: Atomic force microscopy (AFM)	06.05	Dr. Ebru Cihan
6	M3: Low-d materials and application	Carbon Nanostructures	13.05	Dr. Shirong Huang
7	M3: Low-d materials and application	Introduction to low-dimensional materials (LDMs)	20.05	Dr. Shirong Huang
8	M3: Low-d materials and application	Invited talk: Nanostructured materials-based Biosensors	27.05	Dr. Shirong Huang, Prof. Anthony Guiseppi-Elie
9	M3: Low-d materials and application	Nanostructured materials-based Gas sensors	03.06	Dr. Shirong Huang
10	M3: Low-d materials and application	Machine Learning Basics	17.06	Dr. Attila Cangi
11	M3: Low-d materials and application	Artificial Neural Networks	24.06	Dr. Attila Cangi
12	M3: Low-d materials and application	Biomimetic electronic noses (e-noses); AI-powered Sensors	01.07	Dr. Shirong Huang
13	M3: Low-d materials and application	Machine learning for molecule sensing	08.07	Dr. Shirong Huang, Prof. Cuniberti
14	/	Exam day	15.07	Dr. Shirong Huang

- ▶ **Exercise class:** Tomorrow, June 18, 13:00 - 14:30 (MOL/213/H)
- ▶ **Q&A session** on Scientific Project “AI-Driven Prediction of Material Properties”

Why ML for Materials

From AI → ML

Core Vocabulary

Supervised, Unsupervised, Reinforcement Learning

Regression vs. Classification

Loss Function

Linear Regression

Model Evaluation

Under- & Over-fitting

Key Takeaways


The Data Flood in Materials Science

- ▶ High-throughput calculations (first-principles, molecular dynamics, etc.), combinatorial experiments, and open databases have increased significantly
- ▶ Manual analysis or brute-force simulation cannot keep pace
- ▶ Hidden patterns = discovery opportunities → need **Machine Learning**



The Data Flood in Materials Science

- ▶ High-throughput calculations (first-principles, molecular dynamics, etc.), combinatorial experiments, and open databases have increased significantly
- ▶ Manual analysis or brute-force simulation cannot keep pace
- ▶ Hidden patterns = discovery opportunities → need **Machine Learning**




Traditional Modeling vs Machine Learning

Physics-based simulation

- ▶ DFT / MD deliver high fidelity but at high cost: $\mathcal{O}(\text{minutes to hours})$ per structure.


ML surrogate model

- ▶ Train once, then predict new materials in $\mathcal{O}(\text{milliseconds})$.
- ▶ Enables rapid screening and inverse design.



Can We Learn the Mass Density from the Atomic Number?

- ▶ Toy example: atomic number Z vs mass density – relationship is non-linear
- ▶ **Question:** “Could an algorithm automatically learn this trend... and then far more complex ones like band gap or elasticity?”



To answer, we need ML basics: vocabulary, models, and evaluation.

What is AI?



Some trends about AI




<https://hai.stanford.edu/ai-index/2023-ai-index-report>

Some trends about AI

Number of GitHub AI projects, 2011–23


Source: GitHub, 2023 | Chart: 2024 AI Index report



<https://hai.stanford.edu/ai-index/2023-ai-index-report>

Some trends about AI

Number of AI publications in the world, 2010–22
Source: Center for Security and Emerging Technology, 2023 | Chart: 2024 AI Index report




<https://hai.stanford.edu/ai-index/2023-ai-index-report>

Some trends about AI

Number of notable machine learning models by sector, 2003–23

Source: Epoch, 2023 | Chart: 2024 AI Index report




<https://hai.stanford.edu/ai-index/2023-ai-index-report>

Some trends about AI

Foundation models by access type, 2019–23

Source: Bommasani et al., 2023 | Chart: 2024 AI Index report




<https://hai.stanford.edu/ai-index/2023-ai-index-report>

Some trends about AI

Estimated training cost of select AI models, 2017–23


Source: Epoch, 2023 | Chart: 2024 AI Index report



<https://hai.stanford.edu/ai-index/2023-ai-index-report>


Zoom-Out : What Is AI?

- ▶ **Artificial Intelligence (AI)** – umbrella term for algorithms that perform tasks normally requiring human intelligence (planning, perception, reasoning)
- ▶ **Machine Learning (ML)** – *subset* of AI; systems that *learn patterns from data* instead of being hard-coded
- ▶ **Deep Learning (DL)** – *subset* of ML; models based on multi-layer neural networks




Anatomy of Machine Learning

- ▶ Three core **learning paradigms**:
 1. **Supervised** – learn mapping from inputs x to labels y
 2. **Unsupervised** – discover structure in unlabeled data
 3. **Reinforcement** – learn via trial-and-error rewards
- ▶ Each paradigm branches into many *methods* (linear regression, k-means, Q-learning)
- ▶ Common thread: define a goal (loss or reward), optimize it



Where the Paradigms Help in Materials Science

Paradigm	Example materials tasks
Supervised	Predict band gap, elastic modulus, phase label
Unsupervised	Cluster high-entropy alloys; reduce descriptor space with PCA; anomaly detection in synthesis logs
Reinforcement	Active-learning loop in self-driving lab; Bayesian optimization of alloy composition



Property prediction dominates our immediate needs, so we now focus on the **supervised** learning paradigm.

Anatomy of a Dataset

- ▶ A **dataset** = collection of *samples* (rows)
- ▶ Each sample records one material/formula
- ▶ Many columns make up the **features** (descriptors)
- ▶ One column is **label** (target property)

df_12_struct.head()

	structure	log10(K_VRH)	density	vpa	packing fraction	mean ordering parameter shell 1	mean ordering parameter shell 2	mean ordering parameter shell 3
0	[[1.01010394 3.03031182 0.] Li, [3.030...	1.477121	1.825667	17.820044	0.774182	0.411544	0.152153	0.102496
1	[[2.32540495 2.38773782 3.89634095] Al, [0.	1.963788	3.675852	11.659282	0.258695	0.421807	0.181921	0.146386
2	[[0. 0. 0.] Fe, [1.76869807 1.76869807 0. ...	2.307496	8.700535	11.066011	0.958160	0.333333	0.030303	0.014245
3	[[4.91140724 0.82858616 4.10907037] F, [3.1732...	1.301030	5.266007	20.142174	0.886432	0.476271	0.230530	0.181574
4	[[4.68532551 4.68532749 4.68533345] Sc, [1.561...	2.235528	8.631463	15.237572	0.734461	0.551982	0.241225	0.132669

Example snippet: structure, mass density, ordering parameters, packing fraction, bulk modulus


Features & Labels in Materials Science

Features (Descriptors)

- ▶ Numerical representation of composition/structure
- ▶ Examples: average electronegativity, packing fraction, atomic structure descriptor, atomic-site symmetry

Labels (Targets)


- ▶ Property we wish to predict:
Continuous → band gap [eV], formation energy [eV/atom]
Categorical → metal vs insulator, phase (hcp/fcc)



Descriptor (feature) vector feeds model; label guides learning

Model, Parameters & Loss

- ▶ **Model** $\hat{y} = f_{\theta}(x)$ – mathematical function mapping features to label
- ▶ **Parameters** θ – weights the algorithm must learn
- ▶ **Loss** $L(y, \hat{y})$ – quantifies “how wrong” predictions are:
Regression → MSE
Classification → cross-entropy
- ▶ Learning = find $\theta^* = \arg \min_{\theta} \sum L(y, \hat{y})$




Black-box ML model

Feature scaling: Standardize descriptors so weight magnitudes are comparable

Model, Parameters & Loss


- ▶ **Model** $\hat{y} = f_{\theta}(x)$ – mathematical function mapping features to label
- ▶ **Parameters** θ – weights the algorithm must learn
- ▶ **Loss** $L(y, \hat{y})$ – quantifies “how wrong” predictions are:
Regression → MSE
Classification → cross-entropy
- ▶ Learning = find $\theta^* = \arg \min_{\theta} \sum L(y, \hat{y})$



Feature scaling: Standardize descriptors so weight magnitudes are comparable

Train · Validation · Test Split

- ▶ **Training set** – model sees labels; fits weights
- ▶ **Validation set** – monitors generalization; tunes hyper-parameters, triggers early stopping
- ▶ **Test set** – sealed until final evaluation; unbiased estimate of real-world performance
- ▶ Typical split: 70 / 15 / 15 % (flexible)




Golden Rule

Never use the test set during model development.


Three Ways to Learn from Data

- ▶ **Supervised learning**: learn mapping $x \mapsto y$ from labelled data
- ▶ **Unsupervised learning**: find structure/patterns in *unlabelled* data
- ▶ **Reinforcement learning**: learn a policy by maximizing cumulative reward through interaction
- ▶ Common pattern: define objective \Rightarrow optimize it



Supervised Learning in Materials Science


- ▶ Inputs: composition / structure descriptors
- ▶ Targets: *continuous* properties (band gap, elastic modulus) or *categories* (metal vs insulator)
- ▶ Typical algorithms: linear & kernel regression, random forests, neural networks
- ▶ Enables high-throughput **property prediction** and rapid screening



Supervised learning workflow: Descriptor vector → model → predicted property

Unsupervised Learning – When Labels Are Scarce

- ▶ Goal: discover latent structure.
- ▶ Tasks: cluster unknown alloys, reduce high-dimensional descriptor space via PCA, anomaly detection in synthesis logs.
- ▶ Methods: k-means, DBSCAN, autoencoders, PCA, t-SNE.



Reinforcement Learning – When Labels Are Scarce


- ▶ Goal: choose actions to maximize reward.
- ▶ Tasks: active-learning loop in self-driving lab, Bayesian optimization of composition, autonomous furnace tuning.
- ▶ Methods: Q-learning, policy gradients, GP-based Bayesian optimization.

Reinforcement Learning



Which Paradigm Fits Your Problem?

- ▶ **Do you have labels?**
 - ▶ Yes ⇒ **Supervised**.
 - ▶ No ⇒ **Unsupervised**.
- ▶ Need sequential decisions and real-time feedback? ⇒ **Reinforcement**.
- ▶ Example walk-through:
 - ▶ Predict elastic modulus → **Supervised**
 - ▶ Cluster 10 000 unknown materials into categories (metal, insulator, oxides, etc.) → **Unsupervised**
 - ▶ Decide next sample or processing step in autonomous lab → **RL**



Take-away


We'll focus on **supervised learning** next, since it underpins most materials-property prediction tasks.

Two Supervised Question Types

Regression

- ▶ Predict a *continuous* value
- ▶ Example: band gap (eV), elastic modulus (GPa)
- ▶ Loss: MSE / MAE


Regression



Classification


- ▶ Predict a *category label*
- ▶ Example: metal vs insulator, crystal system (hcp/fcc)
- ▶ Loss: cross-entropy

Classification



Why We Need a Loss

- ▶ Model outputs a *prediction* \hat{y} ; reality provides the *true value* y
- ▶ We need a numerical measure of “how wrong” each prediction is
- ▶ **Loss function** $L(y, \hat{y})$ quantifies this error
- ▶ Training = adjust parameters to minimize $\sum L(y, \hat{y})$



Mean-Squared Error (MSE)

- ▶ Standard loss for regression tasks
- ▶ Formula: $\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$
- ▶ Squared term penalizes large errors more strongly
- ▶ Differentiable \Rightarrow gradient-based optimization works smoothly

Loss Functions Overview

- ▶ **Mean-Squared Error (MSE)**

$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$ - Squared term penalizes large errors more strongly

- ▶ **Mean Absolute Error (MAE)** – also for regression and more robust to outliers

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

- ▶ **Root Mean-Squared Error (RMSE)**

$$RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2}$$

- ▶ **Binary Cross Entropy** – for classification

$$BCE = \frac{1}{N} \sum_{i=1}^N [y_i \log \hat{y}_i + (1 - y_i) \log (1 - \hat{y}_i)]$$


- ▶ Choice of loss depends on task and distribution of targets

Key Take-away

Pick a loss that matches your prediction type; for regression we'll typically use MSE.


Linear Model & Assumptions

- ▶ Linear regression is a type of supervised machine-learning algorithm
- ▶ Maps the data points with most optimized linear functions
- ▶ Hypothesis: $\hat{y} = \theta_1 + \theta_2 x$
- ▶ **Assumptions:**
 - ▶ Relationship between input and output is linear
 - ▶ Errors are independent and identically distributed
- ▶ Goal: find θ_1, θ_2 that minimize MSE.



Loss Function in Linear Regression

- ▶ Recall MSE loss
$$L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2$$
- ▶ Calculates the average of the squared errors between the predicted values and the actual values
- ▶ Purpose is to determine optimal values θ_1, θ_2 for the model $\hat{y} = \theta_1 + \theta_2 x$
- ▶ How do we find θ_1, θ_2 ?
- ▶ This is done by minimizing the loss function, for example by gradient descent



Visualizing MSE: Geometric Intuition

- ▶ Each residual $r_i = y_i - \hat{y}_i$ shown as vertical bar
- ▶ Squaring = shaded area of each residual square
- ▶ **Best-fit line** minimizes total shaded area

Visualization

Finding Model Parameters (Closed-form vs. Gradient Descent)

Closed-form (Normal equation):

$$\theta = (\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{y}$$

- ▶ Good for small to moderate dimensionality of features
- ▶ Requires matrix inversion

Gradient Descent (iterative):

$$\theta = \theta - \alpha \cdot \text{"gradient update"}$$

- ▶ Scales to large datasets; underlies neural networks
- ▶ Converges if learning rate α chosen properly

Let's compute the gradient updates specifically for a simple example.

Computing the Gradient Update

- ▶ Consider a simple linear regression model: $\hat{y} = \theta_1 + \theta_2x$
- ▶ Let's compute the gradient update.



Linear regression: Gradient descent

Deriving the gradient update:

$$\theta = \theta - \alpha \cdot \text{"gradient update"}$$

$$= \theta - \alpha \cdot \nabla \theta$$

Consider a simple regression model:

$$\hat{y} = \theta_0 + \theta_1 \cdot x$$

This means we need to calculate how to update θ_0 and θ_1 , so

$$\theta_1 = \theta_1 - \kappa \Delta \theta_1 = \theta_1 - \kappa \frac{\partial L(\theta_0, \theta_1)}{\partial \theta_1}$$

$$\theta_2 = \theta_2 - \kappa \Delta \theta_2 = \theta_2 - \kappa \frac{\partial L(\theta_0, \theta_2)}{\partial \theta_2}$$

Recall that $L(\theta_0, \theta_1) = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$
(MSE loss)

For θ_1 , we then obtain

$$\frac{\partial L}{\partial \theta_1} = \frac{2}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

$$= \frac{1}{N} \sum_{i=1}^N 2(\hat{y}_i - y_i) \frac{\partial}{\partial \theta_1} (\hat{y}_i - y_i)$$

$$= \frac{1}{N} \sum_{i=1}^N 2(\hat{y}_i - y_i) \frac{\partial}{\partial \theta_1} (\theta_0 + \theta_1 \cdot x - y_i)$$

$$= \frac{1}{N} \sum_{i=1}^N 2(\hat{y}_i - y_i) x$$

And for θ_2 , we obtain

$$\frac{\partial L}{\partial \theta_2} = \dots = \frac{1}{N} \sum_{i=1}^N 2(\hat{y}_i - y_i) \frac{\partial}{\partial \theta_2} (\theta_0 + \theta_1 \cdot x - y_i)$$

$$= \frac{1}{N} \sum_{i=1}^N 2(\hat{y}_i - y_i) x$$

we obtain the following gradient updates:

$$\theta_1 = \theta_1 - \kappa \frac{2}{N} \sum_{i=1}^N (\hat{y}_i - y_i) x$$

$$\theta_2 = \theta_2 - \kappa \frac{2}{N} \sum_{i=1}^N (\hat{y}_i - y_i) x$$

Computing the Gradient Update


- We derived the following gradient updates:

- Gradient update for intercept:

$$\theta_1 = \theta_1 - \alpha \frac{\partial L(\theta_1, \theta_2)}{\partial \theta_1} = \theta_1 - \alpha \frac{2}{N} \sum_i^N (\hat{y}_i - y_i)$$

- Gradient update for slope:

$$\theta_2 = \theta_2 - \alpha \frac{\partial L(\theta_1, \theta_2)}{\partial \theta_2} = \theta_2 - \alpha \frac{2}{N} \sum_i^N (\hat{y}_i - y_i) \cdot x_i$$



Beyond Linear Regression

Simple
Linear
Regression

$$y = b_0 + b_1 x_1$$

Multiple
Linear
Regression

$$y = b_0 + b_1 x_1 + b_2 x_2 + \dots + b_n x_n$$

Polynomial
Linear
Regression

$$y = b_0 + b_1 x_1 + b_2 x_1^2 + \dots + b_n x_1^n$$

Regression Metrics

Metric	Formula	Intuition
MSE	$\frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2$	Penalizes large errors (squared)
MAE	$\frac{1}{N} \sum_i^N \hat{y}_i - y_i $	More robust to outliers
R^2	$1 - SS_{res}/SS_{tot}$	Fraction of variance explained (1=perfect)

with


$$SS_{res} = \sum_i^N (\hat{y}_i - y_i)^2$$

$$SS_{tot} = \sum_i^N (\bar{y} - y_i)^2$$

- ▶ Always pair R^2 with an absolute error metric.
- ▶ Choose metric relevant to scientific tolerances.


Predicted vs Actual Plot (Parity Plot)

- ▶ Quick visual sanity check: perfect model lies on 45° line.
- ▶ Systematic bias shows as slope $\neq 1$.
- ▶ Colour/size points by uncertainty or density.




Residual Analysis

- ▶ Plot residual $y - \hat{y}$ vs predicted \hat{y} .
- ▶ Patterns reveal issues:
 - ▶ Funnel shape \Rightarrow non-constant variance.
 - ▶ Curved trend \Rightarrow missing non-linear term.
- ▶ Suggest remedies: log-transform target, add polynomial features, or change model.




Residual Analysis

- ▶ Plot residual $y - \hat{y}$ vs predicted \hat{y} .
- ▶ Patterns reveal issues:
 - ▶ Funnel shape \Rightarrow non-constant variance.
 - ▶ Curved trend \Rightarrow missing non-linear term.
- ▶ Suggest remedies: log-transform target, add polynomial features, or change model.




Model Capacity Spectrum


- ▶ **Under-fit (high bias)**: model too simple, misses trend.
- ▶ **Good fit**: captures true signal, generalizes well.
- ▶ **Over-fit (high variance)**: model too complex, memorizes noise.



High Bias
(Underfitting)




Low Bias, Low Variance
(Goodfitting)



High Variance
(Overfitting)

Model Capacity Spectrum

- ▶ **Under-fit (high bias)**: model too simple, misses trend.
- ▶ **Good fit**: captures true signal, generalizes well.
- ▶ **Over-fit (high variance)**: model too complex, memorizes noise.



Avoiding Overfitting: Lasso and Ridge Regularisation

Lasso (L1)

- ▶ Loss: $L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2 + \lambda \sum_i^N |\theta_i|.$
- ▶ Lasso encourages sparsity by bringing some coefficients to exactly zero.
- ▶ Can set some weights *exactly* zero \Rightarrow implicit feature selection.

Ridge (L2)

- ▶ Loss: $L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2 + \lambda \sum_i^N \theta_i^2$
- ▶ Shrinks all weights smoothly toward zero.
- ▶ Encourages the model to have smaller and more balanced weights.

Elastic net (L1+L2)

- ▶ Loss: $L = \frac{1}{N} \sum_i^N (\hat{y}_i - y_i)^2 + \lambda \sum_i^N |\theta_i| + \lambda \sum_i^N \theta_i^2$
- ▶ Combines both lasso and ridge regularization.

Need to optimize additional hyperparameter λ .

One Slide Summary

- ▶ Core **vocabulary**: dataset, feature, label, model, loss
- ▶ ML **paradigms**: supervised, unsupervised, reinforcement
- ▶ Regression vs. classification; MSE loss; linear-regression baseline
- ▶ Model **evaluation**: train/val/test, metrics, residuals
- ▶ Diagnose **under-/over-fit**; Ridge & Lasso for regularization

